

Syed Ayan Ali
+852 97537220 | syedayanali28@gmail.com
GitHub | LinkedIn

May 4, 2026

Hiring Team
Binance
Hong Kong SAR

Dear Hiring Manager,

I am applying for **Performance Engineer** on the Binance Accelerator Program in Hong Kong SAR. I've worked in HKMA delivery where AI and data hit real production scrutiny and you still own latency after deploy. I want that mindset on large scale systems where performance is the product.

The role blends systems work in Python and Go with observability and disciplined testing before rollout. I stay close to metrics during change and I don't treat model work as a science fair when real users wait on every millisecond.

- **HKMA RAG and observability:** I ship Python services behind RAG on about 10,000 PDFs with SQL vector retrieval and Grafana monitoring for operations plus automation with Next.js and JIRA while I own scrum delivery and L3 support when issues surface.
- **Jay Cee Ess performance and revenue:** I improved performance for end users on a stack where revenue was on the line, cut cart abandonment about 70% after multi wallet checkout work, and cut price lookup latency about 99% with routing using Redis and repeatable integration tests before each release.
- **Universal Innotech data pipelines:** I ran Python ETL across fifty thousand plus records per cycle with validation before each rollout and refactored APIs so client decisions land faster with less brittle handoffs downstream.

I'm strongest in Python, JavaScript, TypeScript, and SQL, with growing depth in Go and in workflows that use LLMs. You can see AI and ML project context on my GitHub link above. Happy to talk on specifics.

Warm Regards,

Syed Ayan Ali